

# Reinforcement Learning

## Introduction and Model-Free Learning

---

Davide Abati

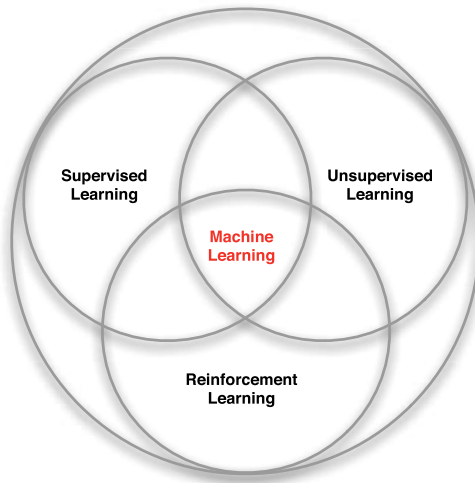
December 9, 2017

University of Modena and Reggio Emilia

All this material is a free re-arrangement of [David Silver's UCL Course on RL](#).  
You are also encouraged to take a look to his [Youtube lectures](#).

**What is reinforcement learning?**

---



What makes reinforcement learning different from other machine learning paradigms?

- There is no supervisor, only a *reward* signal.
- Feedback is delayed, not instantaneous
- Time really matters (sequential, non i.i.d. data)
- Agent is *active*: its actions affect the environment he lives in.

- A **reward**  $R_t$  is a scalar feedback signal
- Indicates how well agent is doing at step  $t$
- The agent's job is to maximise cumulative reward over an episode

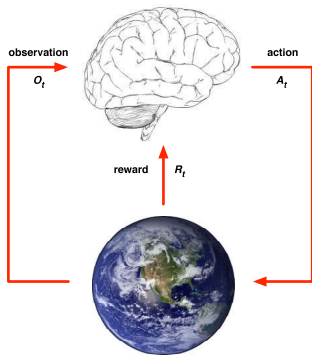
- Fly stunt manoeuvres in a helicopter
  - +ve reward for following desired trajectory
  - -ve reward for crashing
- Defeat the world champion at Go
  - +ve/-ve reward for winning/losing a game
- Make a humanoid robot walk
  - +ve reward for forward motion
  - -ve reward for falling over
- Play Atari games better than humans
  - +ve reward for increasing/decreasing score

## Inside a RL agent

---



- Goal: *select actions to maximise total future reward*
- Actions may have long term consequences
- Reward may be delayed
- It may be better to sacrifice immediate reward to gain more long-term reward.
  - A financial investment may take months to mature
  - Refuelling a helicopter now might prevent a crash in several hours
  - Blocking opponent moves might help winning chances many moves from now



- At each step  $t$  the agent:
  - Receives observation  $O_t$
  - Receives scalar reward  $R_t$
  - Executes action  $A_t$
- The environment:
  - Receives action  $A_t$
  - Emits observation  $O_{t+1}$
  - Emits scalar reward  $R_{t+1}$
- $t$  increments at env. step

- The **history** is the sequence of observations, actions, rewards

$$H_t = O_1, R_1, A_1, \dots, A_{t-1}, O_t, R_t$$

- The **state** is the information used to determine what happens next.
  - It is a function of the history:

$$S_t = f(H_t)$$

## Agent state $S_t^a$

whatever information the agent uses  
to pick the next action

it is the information used by RL algo-  
rithms

## Environment state $S_t^e$

whatever data the environment uses  
to pick the next observation/reward

usually not visible by the agent

- **Full observability**: agent directly observes environment state
- **Partial observability**: agent indirectly observes environment state

- An agent may include one or more of these components:
  - Policy: agent's behaviour function
  - Value function: how good is each state and/or action
  - Model: representation of the environment

- A **policy** is the agent's behaviour
- It is a map from state to action
- Deterministic policy:  $a = \pi(s)$
- Stochastic policy:  $\pi(a|s) = \mathbb{P}[A_t = a | S_t = s]$

- Value function is a prediction of future reward
- Used to evaluate goodness/badness of states
- And therefore to select between actions


$$v_{\pi}(s) = \mathbb{E}_{\pi}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | S_t = s]$$

- A **model** predicts what the environment will do next
- $\mathcal{P}$  predicts the next state
- $\mathcal{R}$  predicts the next (immediate) reward


$$\mathcal{P}^a ss' = \mathbb{P}[S_{t+1} = s' | S_t = s, A_t = a]$$

$$\mathcal{R}_s^a = \mathbb{E}[R_{t+1} | S_t = s, A_t = a]$$






- Rewards: -1 per time-step
- Actions: N, S, W, E
- States: Agent's location




- Arrows represent policy  $\pi(s)$  for each state  $s$


# Maze example: value function



- Numbers represent policy  $v_{\pi}(s)$  for each state  $s$



- Grid layout represent transition model  $\mathcal{P}_{ss'}^a$
- Numbers represent immediate reward  $R_s^a$  from each state  $s$  (same for all  $a$ )



## Model-free prediction

---

- Model-free prediction
  - estimate the value function given a policy in a non-observable environment
    - Monte-Carlo Learning
    - Temporal-Difference Learning

- MC methods learn directly from episodes of experience
- MC is *model-free*: no explicit knowledge of environment mechanisms
- MC learns from complete episodes
  - Caveat: can only apply to *episodic* environments (all episodes must terminate).
- MC uses the simplest possible idea:  $\text{value} = \text{mean return}$



- Goal: learn  $v_\pi$  from episodes of experience under policy  $\pi$

$$S_1, A_1, R_2, \dots, S_k \sim \pi$$

- Recall that the *return* is the total discounted reward:

$$G_t = R_{t+1} + \lambda R_{t+2} + \dots + \lambda^{T-1} R_T$$

- Recall that the value function is the expected return:

$$v_\pi = \mathbb{E}_\pi[G_t | S_t = s]$$

- Monte-Carlo policy evaluation uses *empirical mean* return instead of *expected* return

- To evaluate state  $s$
- Every time-step  $t$  that state  $s$  is visited in an episode,
- Increment counter  $N(s) \leftarrow N(s) + 1$
- Increment total return  $S(s) \leftarrow S(s) + G_t$
- Value is estimated by mean return  $V(s) = S(s)/N(s)$
- By law of large numbers,  $V(s) \rightarrow v_\pi(s)$  as  $N(s) \rightarrow \infty$

- Update  $V(s)$  incrementally after episode  $S_1, A_1, R_2, \dots, S_T$
- Compute return  $G_t$
- For each state  $S_t$  with return  $G_t$

$$N(S_t) \leftarrow N(S_t) + 1$$

$$V(S_t) \leftarrow V(S_t) + \frac{1}{N(S_t)}(G_t - V(S_t))$$

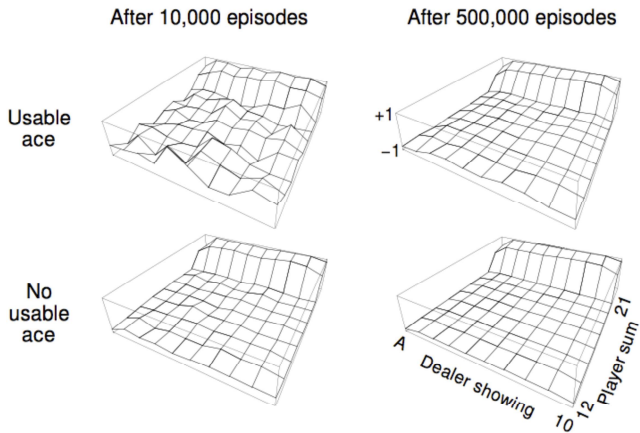
- Usually a running mean is employed, i.e. forget old episodes

$$V(S_t) \leftarrow V(S_t) + \alpha(G_t - V(S_t))$$

- States (200 of them):
  - Current sum (12-21)
  - Dealer's showing card (ace-10)
  - Do I have a useable ace? (yes-no)
- Action **stick**: Stop receiving cards (and terminate)
- Action **twist**: Take another card (no replacement)
- Reward for **stick**:
  - +1 if sum of cards  $>$  sum of dealer cards
  - 0 if sum of cards = sum of dealer cards
  - -1 if sum of cards  $<$  sum of dealer cards
- Reward for **twist**:
  - -1 if sum of cards  $>$  21 (and terminate)
  - 0 otherwise
- Transitions: automatically **twist** if sum of cards  $<$  12



# Blackjack Value Function after Monte-Carlo Learning



Policy: **stick** if sum of cards  $\leq 20$ , otherwise **twist**

- Goal: learn  $v_\pi$  online from experience under policy  $\pi$
- Incremental every-visit Monte-Carlo
  - Update value  $V(S_t)$  toward actual return  $G_t$


$$V(S_t) \leftarrow V(S_t) + \alpha(G_t - V(S_t))$$

- Simplest temporal-difference learning algorithm: TD(0)
  - Update value  $V(S_t)$  toward estimated return  $R_{t+1} + \gamma V(S_{t+1})$


$$V(S_t) \leftarrow V(S_t) + \alpha(R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$

- $R_{t+1} + \gamma V(S_{t+1})$  is called the TD target
- $\delta_t = R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$  is called the TD error

Changes recommended by  
Monte Carlo methods ( $\alpha=1$ )



Changes recommended  
by TD methods ( $\alpha=1$ )



- TD can learn *before* knowing the final outcome
  - TD can learn online after every step
  - MC must wait until end of episode before return is known
- TD can learn without the final outcome
  - TD can learn from incomplete sequences
  - MC can only learn from complete sequences
  - TD works in continuing (non-terminating) environments
  - MC only works for episodic (terminating) environments



- Return  $G_t = R_{t+1} + R_{t+2} + \dots + \gamma^{T-1}R_T$  is *unbiased* estimate of  $v_\pi(S_t)$
- True TD target  $R_{t+1} + v_\pi(S_{t+1})$  is unbiased estimate of  $v_\pi(S_t)$
- TD target  $R_{t+1} + v(S_{t+1})$  is biased estimate of  $v_\pi(S_t)$
- TD target is much lower variance than the return:
  - Return depends on many random actions, transitions, rewards
  - TD target depends on one random action, transition, reward

## Model-free control

---

- Model-free prediction
  - estimate the value function given a policy in a non-observable environment
    - Monte-Carlo Learning
    - Temporal-Difference Learning
- Model-free control
  - find a good policy in a non-observable environment
    - On-Policy Monte-Carlo Control
    - On-Policy Temporal-Difference Learning
    - Off-Policy Learning

- Given a policy  $\pi$
- Evaluate the policy  $\pi$

$$v_{\pi}(s) = \mathbb{E}[R_{t+1} + \gamma R_{t+2} + \dots | S_t = s]$$

$$q_{\pi}(s, a) = \mathbb{E}[R_{t+1} + \gamma R_{t+2} + \dots | S_t = s, A_t = a]$$

- Improve the policy by acting greedily with respect to  $v_{\pi}$

$$\pi' = \text{greedy}(v_{\pi})$$

$$\pi' = \text{greedy}(q_{\pi})$$

- In general, need more iterations of improvement / evaluation
- In model-free contexts, action-value function  $Q(s, a)$  is our only option

# Example of greedy action selection



"Behind one door is tenure - behind the other  
is flipping burgers at McDonald's."


Copyright © 2003 David Farley, d-farley@ibiblio.org

- There are two doors in front of you.
- You open the left door and get reward 0  
 $V(\text{left}) = 0$
- You open the right door and get reward +1  
 $V(\text{right}) = +1$
- You open the right door and get reward +3  
 $V(\text{right}) = +2$
- You open the right door and get reward +2  
 $V(\text{right}) = +2$
- $\vdots$
- Are you sure you've chosen the best door?

- Simplest idea for ensuring continual exploration
- All  $m$  actions are tried with non-zero probability
- With probability  $1 - \epsilon$  choose the greedy action
- With probability  $\epsilon$  choose an action at random

$$\pi(a|s) = \begin{cases} \epsilon/m + 1 - \epsilon, & \text{if } a^* = \arg \max_{a \in \mathcal{A}} Q(s, a) \\ \epsilon/m, & \text{otherwise} \end{cases}$$


- On-policy learning
  - “Learn on the job”
  - Learn about policy  $\pi$  from experience sampled from  $\pi$
- Off-policy learning
  - “Look over someone’s shoulder”
  - Learn about policy  $\pi$  from experience sampled from  $\mu$



**Policy evaluation** Monte-Carlo policy evaluation,  $Q = q_\pi$

**Policy improvement**  $\epsilon$ -greedy policy improvement







Every episode:

Policy evaluation Monte-Carlo policy evaluation,  $Q = q_\pi$


Policy improvement  $\epsilon$ -greedy policy improvement

## Back to the blackjack example






- Temporal-difference (TD) learning has several advantages over Monte-Carlo (MC)
  - Lower variance
  - Online
  - Incomplete sequences
- Natural idea: use TD instead of MC in our control loop
  - Apply TD to  $Q(S, A)$
  - Use  $\epsilon$ -greedy policy improvement
  - Update every time-step



$$Q(S, A) \leftarrow Q(S, A) + \alpha(R + \gamma(Q(S', A') - Q(S, A)))$$




Every time-step:

Policy evaluation with Sarsa,  $Q \approx q_\pi$

Policy improvement with  $\epsilon$ -greedy policy improvement.

```
Initialize  $Q(s, a), \forall s \in S, a \in \mathcal{A}(s)$ , arbitrarily, and  $Q(\text{terminal} - \text{state}, \cdot) = 0$   
for each episode do  
  Intialise  $S$   
  Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)  
  for each step of episode do  
    Take action  $A$ , observe  $R, S'$   
    Choose  $A'$  from  $S'$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)  
     $Q(S, A) \leftarrow Q(S, A) + \alpha(R + \gamma(Q(S', A') - Q(S, A)))$   
     $S \leftarrow S'; A \leftarrow A'$   
  end for  
end for
```




standard  
moves



king's  
moves

- Reward = -1 per time step until reaching goal
- Undiscounted





- Evaluate target policy  $\pi(a, s)$  to compute  $v_\pi(s)$  or  $q_\pi(s, a)$
- While following behaviour policy  $\mu(a|s)$

$$\{S_1, A_1, R_2, \dots, S_T\} \sim \mu$$

- Why is this important?
- Learn from observing humans or other agents
- Re-use experience generated from old policies  $\pi_1, \pi_2, \dots, \pi_{t-1}$
- Learn about *optimal* policy while following *exploratory* policy
- Learn about *multiple* policies while following *one* policy

- We now consider off-policy learning of action-values  $Q(s, a)$
- Next action is chosen using behaviour policy  $A_{t+1} \sim \mu(\cdot|S_t)$
- But we consider alternative successor action  $A' \sim \pi(\cdot|S_t)$
- And update  $Q(S_t, A_t)$  towards value of alternative action

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + Q(S_{t+1}, A') - Q(S_t, A_t))$$

- We now allow both behaviour and target policies to **improve**
- The target policy  $\pi$  is **greedy** w.r.t.  $Q(s, a)$

$$\pi(S_{t+1}) = \arg \max_{a'} Q(S_{t+1}, a')$$

- The behaviour policy  $\mu$  is  **$\epsilon$ -greedy** w.r.t.  $Q(s, a)$
- The Q-learning target then simplifies:

$$\begin{aligned} & R_{t+1} + \gamma Q(S_{t+1}, A') \\ &= R_{t+1} + Q(S_{t+1}, \arg \max_{a'} Q(S_{t+1}, a')) \\ &= R_{t+1} + \max_{a'} Q(S_{t+1}, a') \end{aligned}$$

Initialize  $Q(s, a), \forall s \in S, a \in \mathcal{A}(s)$ , arbitrarily, and  $Q(\text{terminal} - \text{state}, \cdot) = 0$

**for** each episode **do**

    Intialise  $S$

**for** each step of episode **do**

        Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)

        Take action  $A$ , observe  $R, S'$

$Q(S, A) \leftarrow Q(S, A) + \alpha(R + \gamma \max_a (Q(S', a) - Q(S, A)))$

$S \leftarrow S'$

**end for**

**end for**